

Operating Systems Assignment 1– Easy

Name : Jitesh Shamdasani

Entry Number : 2024JCS2043

Login Flow

1. Startup and Console Setup:

- The `init` process starts and ensures that the console device exists and is open.
- Standard input, output, and error are duplicated (via `dup`) so the shell can use them.

2. Infinite Loop and Forking:

- In an infinite loop, `init` forks a child process.
- The parent waits for the child to exit (handling zombie processes).

3. Child Process – Login Procedure:

- In the child (`pid == 0`), the `login()` function is invoked.
- **Login Function Steps:**
 - Prompts the user for a username.
 - Reads input via `gets()` and trims the newline.
 - Compares the input to a defined `USERNAME`.
 - If the username is incorrect, prints an error and increments the attempt counter.
 - If the username is correct, it proceeds to prompt for the password.
 - Reads and trims the password input.
 - Compares the password against the defined `PASSWORD`.
 - If the password matches, the function returns, allowing the login to succeed.
 - After three failed attempts, `login()` delays further execution by calling `sleep()` for a long interval.
- After a successful login, it calls `exec("sh", argv)` to replace the current process image with that of the shell.

4. Outcome:

- Only authenticated users reach the shell. The shell runs in the child process, while the original `init` process continues to monitor and fork new children as needed.

History System Call

1. Global Definitions and Structure

- **Header File:** `history_struct.h`

Defines the `history_struct` and global variables used for history tracking.

- **Structure Definition:**

- `int pid` – Process ID of the process whose history is recorded.
- `char name[16]` – Process name (using 16 chars for consistency with other parts of xv6).
- `int totalMemory` – Memory usage of the process (usually from `proc->sz`).
- `int creationTime` – Increase count at the process creation (for timestamping issues, if used).

- **Global Variables Declared:**

- `hist_arr[MAX_LIMIT]` – Array that holds up to `MAX_LIMIT` history entries.
- `hist_count` – A counter that tracks the number of history entries recorded.
- `hist_lock` – A spinlock to protect updates to the history data.

2. Recording History in the Kernel (in `proc.c`)

- **Initialization:**

In the `pinit(void)` function, the history lock is initialized:

- **Recording a History Entry:**

When a process is created or at appropriate points (likely during fork or process creation), the current process's information is recorded.

3. Retrieving History Data via System Call

In `exec.c` (Actual Implementation)

- **Function:** `gethistory_actual`

- The function acquires `hist_lock` to ensure exclusive access while reading the global history.
- It computes the number of bytes to copy based on `hist_count`.
- Uses `copyout()` to transfer history entries into the user-space buffer.
- Releases the lock and returns the number of history entries.

In `sysproc.c` (System Call Wrapper)

- **Function:** `sys_gethistory`

- `sys_gethistory` obtains the user pointer via `argptr()` and then calls `gethistory_actual()` to do the actual data copy.

- **System Call Number:**

In `syscall.h`, `SYS_gethistory` is defined as `22`, ensuring that when the shell calls `gethistory()`, this system-call gets invoked.

4. Handling the History Command in the Shell (`sh.c`)

- **Detection and Execution of `history` Command:** In the shell's main loop, the code checks if the input command starts with "`history`"
- The shell trims the input.
- It declares a local array to hold the retrieved history.
- Calls the system call `gethistory()`, which internally maps to `sys_gethistory()`.
- Finally, it iterates through the returned history entries and prints them.

6. Overall Flow Summary

1. Process Creation and History Logging:

- During process creation (in `proc.c`), the process's details (PID, name, memory usage) are recorded into `hist_arr` under a lock.

2. User Command for History:

- A user types `history` at the shell prompt.
- The shell recognizes it as a built-in command and calls the system call `gethistory()`.

3. System Call Execution:

- `sys_gethistory()` in `sysproc.c` validates the user's buffer and calls `gethistory_actual()`.
- `gethistory_actual()` in `exec.c` copies the history entries from the kernel's `hist_arr` into the user-provided buffer.

4. Display in Shell:

- The shell prints out the history entries (format: PID, process name, total memory).

5. Block/Unblock (Separate but Similar Flow):

- Block and unblock commands update the per-process `blocked_syscalls` array.
- The kernel checks this array in `syscall()` to enforce blocking behavior before a system call is executed.

This cohesive flow ensures that history tracking and retrieval work seamlessly, while also integrating dynamic control over system call execution via the block/unblock mechanisms.

Block & Unblock

1. User-Space Interface

`usys.S`

The user library macros make the system calls available to user programs.

```
SYSCALL(unblock)
SYSCALL(block)
```

These macros generate the appropriate trap instructions to invoke the kernel functions for block and unblock.

user.h

The user-level prototypes are declared so that user programs (including the shell) can call. This ensures that when the shell's code calls `block(...)` or `unblock(...)`, the call is routed to the corresponding system calls.

2. Kernel System Call Wrappers (sysproc.c)

Block Wrapper

In `sysproc.c`, the `sys_block` function extracts the syscall number from user space and checks that critical syscalls (e.g., fork and exit) are not blocked. It then calls `block_actual` (defined in `exec.c`) to do the actual work.

Unblock Wrapper

Similarly, `sys_unblock` extracts the syscall number and then calls `unblock_actual`:

3. System Call Number Mapping

syscall.h

The syscall numbers for block and unblock are defined so that the kernel's syscall dispatcher can correctly route the calls:

```
#define SYS_block 23
#define SYS_unblock 24
```

syscall.c

The system call array in `syscall.c` maps these numbers to their implementations. When a process invokes the block or unblock system call, the kernel dispatcher calls the corresponding wrapper in `sysproc.c`.

4. Shell Command Parsing (sh.c)

The shell (in `sh.c`) recognizes when the user types a block or unblock command.

This code parses the command line, extracts the target syscall ID, and calls the user-level wrappers, which in turn invoke the kernel system calls.

5. Process Structure and Fork Inheritance (proc.h and proc.c)

In `proc.h`

The `struct proc` is extended with fields for blocking

In `proc.c`

During a fork, the child process inherits the block/unblock settings from the parent

This ensures that the blocking settings are preserved across process creation.

6. Actual Implementation of Block/Unblock (exec.c)

block_actual and unblock_actual Functions

All this perform the actual update of the `blocked_syscalls` array for the current process

these functions ensure that:

- The provided `syscall_id` is within bounds.
 - Critical syscalls (like `fork` and `exit`) are never blocked.
 - The per-process `blocked_syscalls` array is updated correctly.
-

Overall Flow Summary

1. User Interface:

- The shell uses the `block` and `unblock` commands, parsed in `sh.c`.
- The user's command string is converted (using `atoi`) into a syscall number.

2. System Call Invocation:

- The system calls `block()` and `unblock()` are invoked (via `usys.S` and declared in `user.h`).
- These calls enter the kernel and are dispatched through the syscall table (`syscall.h` and `syscall.c`).

3. Kernel Wrappers:

- `sys_block` and `sys_unblock` in `sysproc.c` retrieve the syscall number from user space, perform preliminary checks (e.g., disallowing block of critical syscalls), and then call `block_actual()` or `unblock_actual()`.

4. Actual Implementation:

- In `exec.c`, `block_actual()` and `unblock_actual()` update the current process's `blocked_syscalls` array.
- These changes are inherited by child processes in `fork()` (as defined in `proc.c`).

This complete flow ensures that system calls can be dynamically blocked or unblocked at runtime by updating a per-process array, and the corresponding checks in the syscall dispatcher (in `syscall.c`) will use these settings to either allow or deny the execution of particular system calls.

Chmod System Call

1. User-Space Shell Command (sh.c)

- **Command Parsing:**

- The shell (in `sh.c`) checks if the command begins with "`chmod` " by comparing the first few characters.
- It removes the trailing newline.
- It then skips any extra whitespace and extracts:
 - The target filename (`fileName`).
 - The mode string (`modeStr`).
- The mode string is converted into an integer using `atoi()`.

- **Invocation:**

- The shell then calls the user-level function:

```
if(chmod(fileName, mode) < 0)
    printf(2, "chmod %s failed\n", fileName);
else
    printf(1, "chmod %s succeeded\nnew mode is %d\n", fileName, mode);
```

- This `chmod()` function is declared in `user.h` and invokes the system call via `usys.S`.

2. User-Level System Call Interfaces

- **Prototype in user.h:**

This ensures that when the shell calls `chmod(fileName, mode)`, the call will trap into the kernel.

- **System Call Macro in usys.S:** This macro generates the necessary trap instruction to transfer control to the kernel.

3. Kernel Dispatch: Syscall Mapping (syscall.h and syscall.c)

- **Syscall Number Definition:**

- In `syscall.h`, the syscall number is defined:

```
#define SYS_chmod 25
```

- **Syscall Table Entry:**

- In `syscall.c`, the system call number 25 is mapped to the kernel function `sys_chmod`

```
[SYS_chmod] sys_chmod,
```

4. System Call Wrapper (sysproc.c)

- **Wrapper Function:**
- The function `sys_chmod` acts as a wrapper that extracts the arguments from user space
- It uses `argstr()` and `argint()` to fetch the filename and mode values passed from the user.
- Then, it calls the actual implementation function `chmod_actual(file, mode)`.

5. Actual Implementation (exec.c)

- **Core Functionality:**
- The function `chmod_actual` in `exec.c` performs the permission modification
- **Execution Steps:**
 1. **Validation:**
 - Checks if the target file is "chmod" (to prevent modifying itself).
 - Validates that the mode is within the allowed range `[0, 7]`.
 2. **File System Operation Start:**
 - Calls `begin_op()` to initiate a file system transaction.
 3. **Inode Lookup:**
 - Uses `namei(file)` to obtain a pointer to the inode corresponding to the file.
 - If not found, the operation is terminated.
 4. **Inode Locking/Updating:**
 - Acquires a lock on the inode (`ilock(ip)`).
 - Updates the `mode` field of the inode.
 - Writes changes to disk using `iupdate(ip)` and then unlocks/releases the inode with `iunlockput(ip)`.
 5. **Finalization:**
 - Calls `end_op()` to finish the file system operation.
 - Returns `0` on success or `-1` on any error.

6. Flow Summary

1. **User Action:**
 - The user types `chmod <file> <mode>` in the shell.
2. **Shell Processing (sh.c):**
 - Parses the command to extract `fileName` and `mode`.
 - Invokes the user-level `chmod(fileName, mode)` function.
3. **User-Level to Kernel Transition:**

- The system call interface (via `usys.S`) traps into the kernel.

4. Kernel Dispatch:

- The syscall number `SYS_chmod` (25) is matched and routed to `sys_chmod` (in `sysproc.c`).

5. Wrapper Function (`sysproc.c`):

- `sys_chmod` extracts the arguments and calls `chmod_actual(file, mode)`.

6. Actual File Operation (`exec.c`):

- `chmod_actual` validates inputs, locates the file's inode, updates its mode, and finalizes the file system transaction.

7. Result:

- On success, the new file mode is set, and a success message is printed by the shell. Otherwise, an error message is displayed.
-